

INFO411 Data Mining and Knowledge Discovery

Assignment 2

Audric Ng Si Kai

7431168

```
> tree.cw.train
node), split, n, deviance, yval, (yprob)
      * denotes terminal node

1) root 981 2021.000 2 ( 0.23038 0.51274 0.25688 )
  2) functionary < 0.5 709 1344.000 2 ( 0.13540 0.57546 0.28914 )
    4) FI30.credit.score < 0.5 58 61.720 3 ( 0.00000 0.22414 0.77586 )
      *
        5) FI30.credit.score > 0.5 651 1211.000 2 ( 0.14747 0.60676 0.24578 )
          10) re.balanced..paid.back..a.recently.overdrawn.current.account < 0.
5 57 98.140 3 ( 0.07018 0.33333 0.59649 ) *
          11) re.balanced..paid.back..a.recently.overdrawn.current.account > 0.
5 594 1078.000 2 ( 0.15488 0.63300 0.21212 ) *
        3) functionary > 0.5 272 556.800 1 ( 0.47794 0.34926 0.17279 )
          6) re.balanced..paid.back..a.recently.overdrawn.current.account < 0.5
11 12.890 3 ( 0.00000 0.27273 0.72727 ) *
          7) re.balanced..paid.back..a.recently.overdrawn.current.account > 0.5
261 521.400 1 ( 0.49808 0.35249 0.14943 )
            14) FI30.credit.score < 0.5 9 9.535 3 ( 0.00000 0.22222 0.77778 )
              *
            15) FI30.credit.score > 0.5 252 489.500 1 ( 0.51587 0.35714 0.12698 )
              *
            ) *
```

b)

I created a new dataframe to represent the data of a hypothetical customer. First, I created a new dataframe for the customer called medianCust and created a new dataset with random values to simulate the records of a new customer. Then, I bind the new dataset with the new dataframe created. And lastly, I added the names of the columns of the original dataset to the new dataset for medianCust.

```
> cust.pred = predict(tree.cw.train, medianCust, type = "class")
> cust.pred
[1] 2
```

Levels: 1 2 3

c)

```
> confusion = with(cw.test, table(tree.pred, credit.rating))
> confusion
      credit.rating
tree.pred  1    2    3
      1 162   85   37
      2  90  361  143
      3   5   21   77

> sum(diag(confusion))/sum(confusion)*100
[1] 61.16208
```

Total value = 162 + 85 + 37 + 90 + 361 + 143 + 5 + 21 + 77 = 981

Diag = 162 + 361 + 77 = 600

Accuracy rate = 600/981 x 100 = 61.16%

From the confusion matrix, the decision tree model has a moderately high accuracy for its prediction.

d)

```
> beforeCountFreq = table(cw.train$credit.rating)
> beforeCountFreq
      1    2    3
226 503 252

> # Find the probability of each class
> beforeClassProb = beforeCountFreq/sum(beforeCountFreq)
> beforeClassProb
      1      2      3
0.2303772 0.5127421 0.2568807

> # calculate entropy (before split)
> beforeEntropy = -sum(beforeClassProb * log2(beforeClassProb))
> beforeEntropy
[1] 1.485749

> # First split on functionary > 0.5
> # functionary == 0
> countFreq0 = table(cw.train$credit.rating[cw.train$functionary == 0])
> countFreq0
      1    2    3
 96 408 205

> # Find the probability of each class
> classProb0 = countFreq0/sum(countFreq0)
> (functionaryEnt0 = -sum(classProb0 * log2(classProb0)))
[1] 1.366963
> functionaryEnt0
[1] 1.366963

> # functionary == 1
> countFreq1 = table(cw.train$credit.rating[cw.train$functionary == 1])
> countFreq1
      1    2    3
 96 408 205

> # Find the probability of each class
> classProb1 = countFreq1/sum(countFreq1)
> classProb1
```

```

      1      2      3
0.1354020 0.5754584 0.2891396
> (functionaryEnt1 = -sum(classProb1 * log2(classProb1)))
[1] 1.366963
> functionaryEnt1
[1] 1.366963
> # Numerical value of the gain in entropy
> ent = (beforeEntropy - (functionaryEnt0 * sum(countFreq0) + functionaryEnt1 * sum(countFreq1))/sum(sum(countFreq0) + sum(countFreq1)))
> ent
[1] 0.1187861

```

e)

```

> rf.cw.train = randomForest(credit.rating~., data = cw.train)
> rf.cw.train

```

Call:

```

randomForest(formula = credit.rating ~ ., data = cw.train)
      Type of random forest: classification
      Number of trees: 500
No. of variables tried at each split: 6

```

OOB estimate of error rate: 42.61%

Confusion matrix:

	1	2	3	class.error
1	51	175	0	0.7743363
2	32	448	23	0.1093439
3	14	174	64	0.7460317

```

> rf.pred = predict(rf.cw.train, cw.test[,-46])

```

Total value = 51 + 175 + 0 + 32 + 448 + 23 + 14 + 174 + 64 = 981

Diag = 51 + 448 + 64 = 508

Accuracy rate = 508/981 x 100 = 51.78%

Using the random forest model to train the training dataset, the accuracy rate is lower than the accuracy rate of the decision tree test set. Therefore, there is no improvement in the accuracy rate.

f)

```

> confusionRF = with(cw.test, table(rf.pred, credit.rating))
> confusionRF

```

	credit.rating		
rf.pred	1	2	3
1	55	42	17
2	200	411	194
3	2	14	46

```

> # Question 2f: Overall accuracy
> sum(diag(confusionRF))/sum(confusionRF)*100
[1] 52.19164

```

Total value = 55 + 42 + 17 + 200 + 411 + 194 + 2 + 14 + 46 = 981

Diag = 55 + 411 + 46 = 512

Accuracy rate = 512/981 x 100 = 52.19%

From the confusion matrix, the random forest model has a moderately high accuracy, but it is lower than the accuracy of the decision tree model.

3a)

```

> predict(svmfit, medianCust, decision.values = TRUE)

```

```

1
2

```

```
attr(,"decision.values")
      2/1      2/3      1/3
1 1.021296 1.511396 -0.04938262
Levels: 1 2 3
```

b)

```
> svm.pred = predict(svmfit, cw.test[,-46])
> confusionSVM = with(cw.test, table(svm.pred, credit.rating))
> confusionSVM
      credit.rating
svm.pred 1    2    3
1    109   56   22
2    143  393  162
3     5    18   73
```

Total value = 109 + 56 + 22 + 143 + 393 + 162 + 5 + 18 + 73 = 981

Diag = 109 + 393 + 73 = 575

Accuracy rate = 575/981 x 100 = 58.60%

From the confusion matrix, the SVM(support vector machine) model has a moderately high accuracy, but it is lower than the accuracy of the decision tree model.

c)

```
> summary(tune.svm(credit.rating~., data = cw.train,
+                  kernel = "radial", cost = 10^c(0:2),
+                  gamma = 10^c(-4:-1)))
```

Parameter tuning of 'svm':

- sampling method: 10-fold cross validation

- best parameters:

```
gamma cost
0.001  100
```

- best performance: 0.3965986

- Detailed performance results:

	gamma	cost	error	dispersion
1	1e-04	1	0.4872501	0.04769583
2	1e-03	1	0.4811276	0.04639319
3	1e-02	1	0.3986807	0.06673036
4	1e-01	1	0.4872501	0.04769583
5	1e-04	10	0.4658627	0.04869084
6	1e-03	10	0.3966296	0.05072210
7	1e-02	10	0.4659658	0.06918890
8	1e-01	10	0.4872501	0.04769583
9	1e-04	100	0.3976500	0.05108128
10	1e-03	100	0.3965986	0.05235591
11	1e-02	100	0.4781591	0.06583392
12	1e-01	100	0.4872501	0.04769583

```
> # Fit a model using SVM
> svmTuned = svm(credit.rating~., data = cw.train,
+               kernel = "radial", cost = 10,
+               gamma = 0.001)
> print(svmTuned)
```

Call:

```
svm(formula = credit.rating ~ ., data = cw.train, kernel = "radial", cost
= 10,
    gamma = 0.001)
```

Parameters:
 SVM-Type: C-classification
 SVM-Kernel: radial
 cost: 10

Number of Support Vectors: 834

```
> # Predict the values on test set
> svmTuned.pred = predict(svmTuned, cw.test[,-46])
> # Produce confusion matrix
> confusionTunedSVM = with(cw.test, table(credit.rating, svmTuned.pred))
> confusionTunedSVM
```

	svmTuned.pred		
credit.rating	1	2	3
1	160	92	5
2	87	361	19
3	39	147	71

```
> # Overall accuracy rate
> sum(diag(confusionTunedSVM))/sum(confusionTunedSVM)*100
[1] 60.34659
```

Total value = 160 + 92 + 5 + 87 + 361 + 19 + 39 + 147 + 71 = 981

Diag = 160 + 361 + 71 = 592

Accuracy rate = 592/981 x 100 = 60.35%

From the confusion matrix, the SVM model accuracy after tuning increased by approximately 1.75%, but it is still lower than the accuracy of the decision tree model.

4a)

```
> nb = naiveBayes(credit.rating~., data = cw.train)
> predict(nb, medianCust, type = 'class')
[1] 1
Levels: 1 2 3
> predict(nb, medianCust, type = 'raw')
      1      2      3
[1,] 0.9850729 0.01393277 0.0009942948
> # predict the values on test set
> nb.pred = predict(nb, cw.test[,-46])
> # produce confusion matrix
> confusionNB = with(cw.test, table(nb.pred, credit.rating))
> confusionNB
```

	credit.rating		
nb.pred	1	2	3
1	252	439	173
2	0	4	6
3	5	24	78

```
> # calculate the accuracy rate
> sum(diag(confusionNB))/sum(confusionNB)*100
[1] 34.04689
```

b)

```
> nb
```

Naive Bayes Classifier for Discrete Predictors

Call:

```
naiveBayes.default(x = X, y = Y, laplace = laplace)
```

A-priori probabilities:

Y	1	2	3
---	---	---	---

0.2303772 0.5127421 0.2568807

Conditional probabilities:

functionary

Y	[,1]	[,2]
1	0.5752212	0.4954066
2	0.1888668	0.3917924
3	0.1865079	0.3902912

re.balanced..paid.back..a.recently.overdrawn.current.account

Y	[,1]	[,2]
1	0.9823009	0.1321481
2	0.9542744	0.2090974
3	0.8095238	0.3934582

FI30.credit.score

Y	[,1]	[,2]
1	1.0000000	0.0000000
2	0.9701789	0.1702628
3	0.7936508	0.4054894

gender

Y	[,1]	[,2]
1	0.5265487	0.5004030
2	0.4015905	0.4907079
3	0.3531746	0.4789075

x0..accounts.at.other.banks

Y	[,1]	[,2]
1	2.898230	1.370579
2	3.079523	1.410560
3	3.047619	1.433004

credit.refused.in.past.

Y	[,1]	[,2]
1	0.05752212	0.2333544
2	0.09940358	0.2995010
3	0.21428571	0.4111425

years.employed

Y	[,1]	[,2]
1	3.013274	1.409429
2	2.972167	1.412530
3	3.039683	1.314345

savings.on.other.accounts

Y	[,1]	[,2]
1	3.442478	1.854427
2	3.626243	1.802905
3	3.420635	1.748548

self.employed.

Y	[,1]	[,2]
1	0.1637168	0.3708398
2	0.2206759	0.4151152
3	0.2142857	0.4111425

max..account.balance.12.months.ago

Y	[,1]	[,2]
1	2.955752	1.453819
2	2.978131	1.425968
3	2.940476	1.408719

min..account.balance.12.months.ago

Y	[,1]	[,2]
1	2.973451	1.391787
2	2.956262	1.412126
3	3.115079	1.393872

```

    avrg..account.balance.12.months.ago
Y      [,1]      [,2]
1 3.225664 1.450657
2 2.982107 1.366094
3 3.015873 1.414124

```

```

    max..account.balance.11.months.ago
Y      [,1]      [,2]
1 2.884956 1.390458
2 2.992048 1.412782
3 3.059524 1.436723

```

```

    min..account.balance.11.months.ago
Y      [,1]      [,2]
1 2.827434 1.442636
2 2.990060 1.370543
3 2.984127 1.405647

```

```

    avrg..account.balance.11.months.ago
Y      [,1]      [,2]
1 3.017699 1.391928
2 2.978131 1.361654
3 2.996032 1.407147

```

```

    max..account.balance.10.months.ago
Y      [,1]      [,2]
1 3.026549 1.454252
2 3.001988 1.445558
3 2.968254 1.413856

```

```

    min..account.balance.10.months.ago
Y      [,1]      [,2]
1 2.792035 1.419273
2 3.031809 1.434833
3 2.904762 1.416626

```

```

    avrg..account.balance.10.months.ago
Y      [,1]      [,2]
1 2.946903 1.325582
2 2.998012 1.410686
3 3.027778 1.470569

```

```

    max..account.balance.9.months.ago
Y      [,1]      [,2]
1 3.110619 1.467022
2 2.912525 1.414320
3 2.980159 1.440591

```

```

    min..account.balance.9.months.ago
Y      [,1]      [,2]
1 2.920354 1.363926
2 2.914513 1.419362
3 3.067460 1.385545

```

5a)

The decision tree model looks to be the best as the test set has the highest accuracy rate of 61.16%. This is as compared to the other models, random forest with 52.19%, SVM with 60.35%, and naive bayes with 34.05%.

b)

The Naïve Bayes classifier predicted the conditional probability of the first A-priori probability using the FI3o credit score as a zero. This means that the classifier may have encountered some issues in predicting the A-priori probability using the FI3o credit score.

6a)

```
> glm.fit <- glm((credit.rating==1)~., data = cw.train, family = binomial)
> options(width = 130)
```

b)

```
> summary(glm.fit)
```

```
Call:
glm(formula = (credit.rating == 1) ~ ., family = binomial, data = cw.train)
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-2.00215	-0.65353	-0.42668	-0.00012	2.70789

Coefficients:

	Estimate	Std
. Error z value Pr(> z)		
(Intercept)	-17.551605	429
.995589 -0.041 0.96744		
functionary	1.740533	0
.183036 9.509 < 2e-16 ***		
re.balanced..paid.back..a.recently.overdrawn.current.acount	1.501222	0
.550965 2.725 0.00644 **		
FI30.credit.score	16.502759	429
.993845 0.038 0.96939		
gender	0.577104	0
.178807 3.228 0.00125 **		
x0..accounts.at.other.banks	-0.027413	0
.063141 -0.434 0.66417		
credit.refused.in.past.	-0.935877	0
.341848 -2.738 0.00619 **		
years.employed	0.672572	0
.269126 2.499 0.01245 *		
savings.on.other.accounts	-0.548195	0
.204670 -2.678 0.00740 **		
self.employed.	-0.376394	0
.236506 -1.591 0.11150		
max..account.balance.12.months.ago	-0.004444	0
.062647 -0.071 0.94345		
min..account.balance.12.months.ago	0.030192	0
.063737 0.474 0.63572		
avrg..account.balance.12.months.ago	0.124651	0
.065028 1.917 0.05525 .		
max..account.balance.11.months.ago	-0.010150	0
.063924 -0.159 0.87385		
min..account.balance.11.months.ago	-0.110469	0
.064328 -1.717 0.08593 .		
avrg..account.balance.11.months.ago	0.052783	0
.065196 0.810 0.41816		
max..account.balance.10.months.ago	0.019305	0
.062526 0.309 0.75750		
min..account.balance.10.months.ago	-0.101696	0
.063199 -1.609 0.10759		
avrg..account.balance.10.months.ago	-0.050933	0
.065720 -0.775 0.43834		
max..account.balance.9.months.ago	0.096730	0
.062586 1.546 0.12221		
min..account.balance.9.months.ago	-0.038009	0
.064765 -0.587 0.55728		
avrg..account.balance.9.months.ago	-0.032928	0
.062640 -0.526 0.59912		
max..account.balance.8.months.ago	-0.019017	0
.063459 -0.300 0.76443		

min..account.balance.8.months.ago	-0.041455	0
.062710 -0.661 0.50858		
avrg..account.balance.8.months.ago	-0.106852	0
.063685 -1.678 0.09338		
max..account.balance.7.months.ago	-0.018414	0
.063321 -0.291 0.77120		
min..account.balance.7.months.ago	-0.094176	0
.063702 -1.478 0.13930		
avrg..account.balance.7.months.ago	-0.074021	0
.061950 -1.195 0.23215		
max..account.balance.6.months.ago	0.069171	0
.064686 1.069 0.28492		
min..account.balance.6.months.ago	-0.033830	0
.062428 -0.542 0.58788		
avrg..account.balance.6.months.ago	-0.025278	0
.062786 -0.403 0.68724		
max..account.balance.5.months.ago	0.015218	0
.061902 0.246 0.80581		
min..account.balance.5.months.ago	-0.088221	0
.064391 -1.370 0.17066		
avrg..account.balance.5.months.ago	-0.072089	0
.063401 -1.137 0.25553		
max..account.balance.4.months.ago	0.034718	0
.062889 0.552 0.58091		
min..account.balance.4.months.ago	-0.036728	0
.064179 -0.572 0.56714		
avrg..account.balance.4.months.ago	0.020068	0
.063954 0.314 0.75368		
max..account.balance.3.months.ago	-0.144584	0
.062966 -2.296 0.02166 *		
min..account.balance.3.months.ago	0.014149	0
.064191 0.220 0.82554		
avrg..account.balance.3.months.ago	-0.010770	0
.064635 -0.167 0.86767		
max..account.balance.2.months.ago	0.100711	0
.063196 1.594 0.11102		
min..account.balance.2.months.ago	-0.065585	0
.063059 -1.040 0.29832		
avrg..account.balance.2.months.ago	-0.038225	0
.064392 -0.594 0.55276		
max..account.balance.1.months.ago	-0.073012	0
.065482 -1.115 0.26486		
min..account.balance.1.months.ago	-0.000658	0
.062229 -0.011 0.99156		
avrg..account.balance.1.months.ago	-0.068570	0
.064302 -1.066 0.28626		

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 1058.95 on 980 degrees of freedom
Residual deviance: 820.79 on 935 degrees of freedom
AIC: 912.79

Number of Fisher Scoring iterations: 16

c)

The coefficient for FI3o credit score is abnormally high at 16.50 as compared to the intercept value of -17.55. The coefficient also does not fall in the same range as the other attributes of the dataset from -1 to 1.

d)

```
> # Fit an svm model of your choice to the training set
> summary(tune.svm((credit.rating==1)~., data = cw.train,
+                 kernel = "radial", cost = 10^c(-2:2),
+                 gamma = 10^c(-4:1), type = "C"))
```

Parameter tuning of 'svm':

- sampling method: 10-fold cross validation

- best parameters:

gamma	cost
0.001	100

- best performance: 0.2171408

- Detailed performance results:

	gamma	cost	error	dispersion
1	1e-04	1e-02	0.2303649	0.02447423
2	1e-03	1e-02	0.2303649	0.02447423
3	1e-02	1e-02	0.2303649	0.02447423
4	1e-01	1e-02	0.2303649	0.02447423
5	1e+00	1e-02	0.2303649	0.02447423
6	1e+01	1e-02	0.2303649	0.02447423
7	1e-04	1e-01	0.2303649	0.02447423
8	1e-03	1e-01	0.2303649	0.02447423
9	1e-02	1e-01	0.2303649	0.02447423
10	1e-01	1e-01	0.2303649	0.02447423
11	1e+00	1e-01	0.2303649	0.02447423
12	1e+01	1e-01	0.2303649	0.02447423
13	1e-04	1e+00	0.2303649	0.02447423
14	1e-03	1e+00	0.2303649	0.02447423
15	1e-02	1e+00	0.2303649	0.02250408
16	1e-01	1e+00	0.2303649	0.02447423
17	1e+00	1e+00	0.2303649	0.02447423
18	1e+01	1e+00	0.2303649	0.02447423
19	1e-04	1e+01	0.2303649	0.02447423
20	1e-03	1e+01	0.2293445	0.03000640
21	1e-02	1e+01	0.2303649	0.03628059
22	1e-01	1e+01	0.2303649	0.02447423
23	1e+00	1e+01	0.2303649	0.02447423
24	1e+01	1e+01	0.2303649	0.02447423
25	1e-04	1e+02	0.2293445	0.03000640
26	1e-03	1e+02	0.2171408	0.02372726
27	1e-02	1e+02	0.2395382	0.05200855
28	1e-01	1e+02	0.2303649	0.02447423
29	1e+00	1e+02	0.2303649	0.02447423
30	1e+01	1e+02	0.2303649	0.02447423

```
> (svm2 = svm((credit.rating==1)~., data = cw.train, type = "C"))
```

Call:

```
svm(formula = (credit.rating == 1) ~ ., data = cw.train, type = "C")
```

Parameters:

SVM-Type:	C-classification
SVM-Kernel:	radial
cost:	1

Number of Support Vectors: 664

```
> #Fit a model using svm
> svm.fit = svm((credit.rating==1) ~ ., data = cw.train, type = 'C',
+               gamma = 0.001, cost = 100, kernel = "radial")
> print(svm.fit)
```

Call:

```
svm(formula = (credit.rating == 1) ~ ., data = cw.train, type = "C", gamma  
= 0.001, cost = 100, kernel = "radial")
```

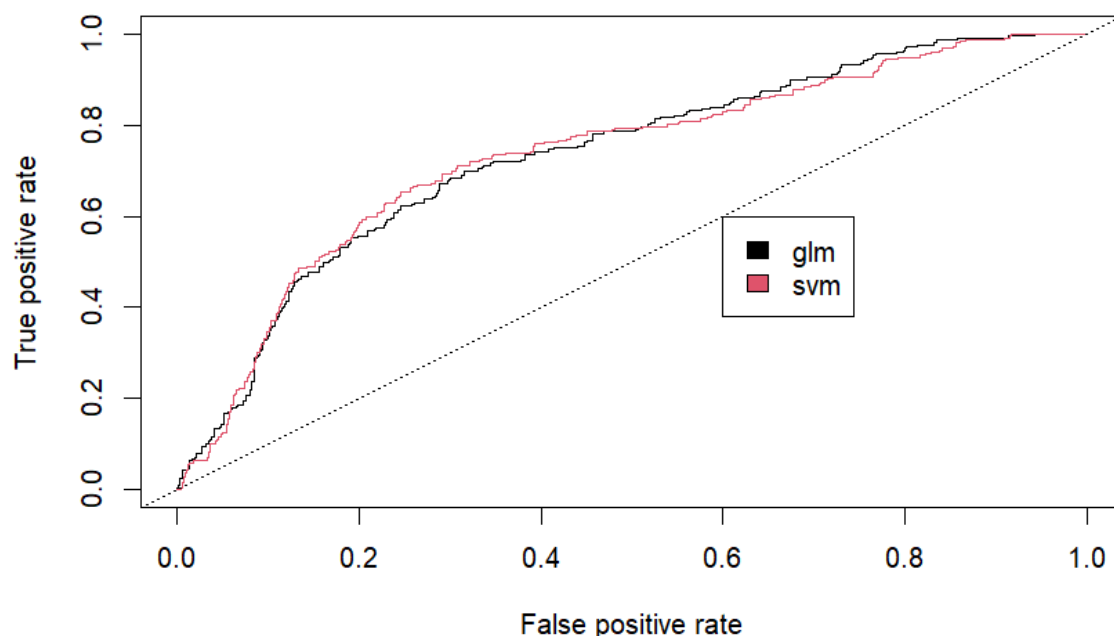
Parameters:

```
SVM-Type: C-classification  
SVM-Kernel: radial  
cost: 100
```

Number of Support Vectors: 522

e)

```
> # Predict the values on test set[SVM]  
> svm.fit.pred = predict(svm.fit, cw.test[,-46], decision.values = TRUE)  
> # Predict the values on test set[GLM]  
> glm.fit.pred = predict(glm.fit, cw.test[,-46])  
> # Make prediction using SVM  
> confusionSVM = prediction(-attr(svm.fit.pred, "decision.values"),  
+                             cw.test$credit.rating == 1)  
> # Create rocs curve based on prediction  
> rocSVM <- performance(confusionSVM, "tpr", "fpr")  
> # Make prediction using Logistic Regression  
> confusionGLM = prediction(glm.fit.pred, cw.test$credit.rating == 1)  
> # create rocs curve based on prediction  
> rocGLM <- performance(confusionGLM, "tpr", "fpr")  
> # Plot the graph  
> plot(rocGLM, col = 1)  
> plot(rocSVM, col = 2, add = TRUE)  
> abline(0, 1, lty=3)  
> # Add the legend to the graph  
> legend(0.6, 0.6, c('glm','svm'), 1:2)
```



Even though both models started at the same false positive rate, the SVM model has a higher false positive rate than the GLM model. Whereas, both models have the same true positive rate at the end of the prediction.